

**Experiment - 1**

**Aim-**To predict Bicarbonate(ppm) present in the well water of Northwest Texas data via Linear Regression Machine learning model.

```
cd /content/drive/MyDrive/Machine Learning/Colab Notebooks/ML
Practicals/1_Practical/Linear regression P1
```

```
/content/drive/MyDrive/Machine Learning/Colab Notebooks/ML
Practicals/1_Practical/Linear regression P1
```

```
ls
```

```
edcCO2.csv      'Ground Water Survey.csv'
fruitohms.csv   'Linear regression_1.ipynb'
```

Importing Required Libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
#import data set
dataset = pd.read_csv('Ground Water Survey.csv')
X= dataset.iloc[:, :-1].values
Y= dataset.iloc[:, 1].values
dataset.head()
```

```
      X      Y
0  7.6  157
1  7.1  174
2  8.2  175
3  7.5  188
4  7.4  171
```

In the following data

X = pH of well water

Y = Bicarbonate (parts per million) of well water

The data is by water well from a random sample of wells in Northwest Texas. Reference: Union Carbide Technical Report K/UR-1

```
dataset.tail()
```

```
      X      Y
29  8.5   48
30  7.8  147
31  6.7  117
```

```
32  7.1  182
33  7.3   87
```

Bicarbonate can be found in water with a pH between 4.3 and 12.3. Above a pH of 8.3, carbonate is also present.

*#Splitting the data*

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test= train_test_split(X,Y,test_size= 0.7)
```

*#Fitting Simple Linear Regression ipynb*

*#This is called Model*

```
from sklearn.linear_model import LinearRegression
regressor= LinearRegression()
regressor.fit(X_train,Y_train)
```

```
LinearRegression()
```

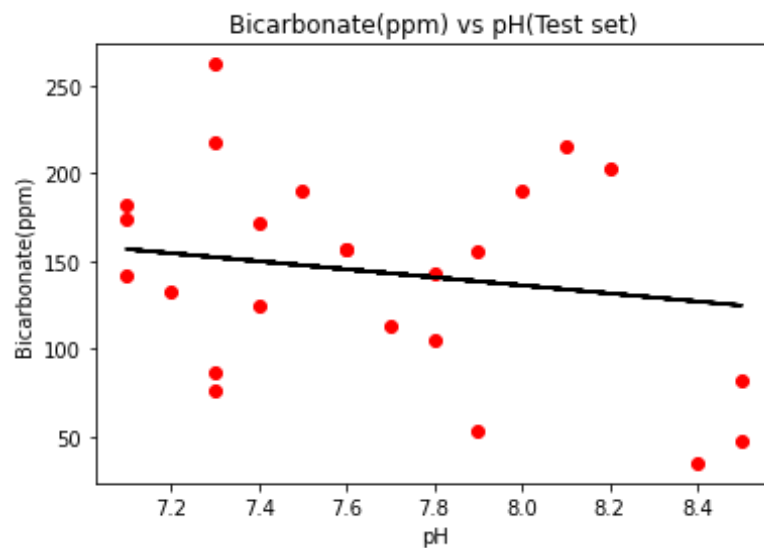

*##Predicting the test results*

```
Y_pred= regressor.predict(X_test)
```

*#Visualising the training set Results*

```
plt.scatter(X_train, Y_train, color='Purple')
plt.plot(X_train, regressor.predict(X_train), color='black')
plt.title('Bicarbonate(ppm) vs pH(Training set)')
plt.xlabel('pH')
plt.ylabel('Bicarbonate(ppm)')
plt.show()
```

```
plt.scatter(X_test, Y_test, color='red')
plt.plot(X_test, regressor.predict(X_test), color='black')
plt.title('Bicarbonate(ppm) vs pH(Test set)')
plt.xlabel('pH')
plt.ylabel('Bicarbonate(ppm)')
plt.show()
```



```
print(regressor.predict([[7.6]]))
```

```
[145.2435247]
```

**Now we will perform the prediction of Bicarbonate(ppm) Present in the well water.**

```
a=float(input("What is the pH of your well water? "))
print('The Bicarbonate (parts per million) in your well water',
regressor.predict([[a]]))
```

```
What is the pH of your well water? 7.1
```

```
The Bicarbonate (parts per million) in your well water [156.6787717]
```

**Conclusion-** Hence we are able to predict the Bicarbonate(ppm) present in the well water of Northeast Texas by training the Linear Regression model with the Water Survey Dataset.

## Experiment - 2

**Aim-** To predict Coronary Heart Disease using Logistic Regression Classifier

**Logistic Regression:** The target variable has three or more nominal categories such as predicting the type of Wine. Ordinal Logistic Regression: the target variable has three or more ordinal categories such as restaurant or product rating from 1 to 5. Model building in Scikit-learn Let's build the diabetes prediction model.

Here, We are going to predict Coronary Heart Disease using Logistic Regression Classifier.

Let's first load the required Coronary Heart Disease dataset using the pandas' read CSV function.

We will download data from the following link:

<https://www.kaggle.com/datasets/billbasener/coronary-heart-disease?resource=download>

```
from google.colab import drive
```

```
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).

```
cd /content/gdrive/MyDrive/Machine Learning/Colab  
Notebooks/ML_Practicals/1_Practical/Logistic regression P2
```

```
/content/gdrive/MyDrive/Machine Learning/Colab  
Notebooks/ML_Practicals/1_Practical/Logistic regression P2
```

```
ls
```

```
CHDdata.csv  CHD_Data.csv  CHDdata.gsheet  Logistic_Regression.ipynb
```

```
import pandas as pd  
col_names = ['Systolic BP', 'Tobacco', 'low-density lipoprotein',  
'Adiposity', 'Famhist', 'typea', 'Obesity', 'Alcohol', 'Age', 'Chd']  
# Load dataset  
CHD = pd.read_csv("CHD_Data.csv", header=None, names=col_names)
```

### **Context**

The data set CHDdata.csv contains cases of coronary heart disease (CHD) and variables associated with the patient's condition: systolic blood pressure, yearly tobacco use (in kg), low density lipoprotein (ldl), adiposity, family history (0 or 1), type A personality score (typea), obesity (body mass index), alcohol use, age, and the diagnosis of CHD (0 or 1).

CHD.head()

	Systolic BP	Tobacco	low-density lipoprotein	Adiposity	Famhist	typea
1	160	12.00	5.73	23.11	Present	49
2	144	0.01	4.41	28.61	Absent	55
3	118	0.08	3.48	32.28	Present	52
4	170	7.50	6.41	38.03	Present	51
5	134	13.60	3.50	27.78	Present	60

	Obesity	Alcohol	Age	Chd
1	25.30	97.20	52	1
2	28.87	2.06	63	1
3	29.14	3.81	46	0
4	31.99	24.26	58	1
5	25.99	57.34	49	1

CHD.tail()

	Systolic BP	Tobacco	low-density lipoprotein	Adiposity	Famhist	typea
458	214	0.4	5.98	31.72	Absent	64
459	182	4.2	4.41	32.10	Absent	52
460	108	3.0	1.59	15.23	Absent	40
461	118	5.4	11.61	30.79	Absent	64
462	132	0.0	4.82	33.41	Present	62

	Obesity	Alcohol	Age	Chd
458	28.45	0.00	58	0
459	28.61	18.72	52	1
460	20.09	26.64	55	0
461	27.35	23.97	40	0
462	14.70	0.00	46	1

CHD.drop('Famhist', inplace=True, axis=1)

**Selecting Feature** Here, we need to divide the given columns into two types of variables dependent(or target variable) and independent variable(or feature variables).

*#split dataset in features and target variable*

```
feature_cols = ['Systolic BP', 'Tobacco', 'low-density lipoprotein',
                'Adiposity', 'typea', 'Obesity', 'Alcohol', 'Age', ]
```

```
X = CHD[feature_cols] # Features
```

```
y = CHD.Chd # Target variable
```

**Splitting Data** To understand model performance, dividing the dataset into a training set and a test set is a good strategy.

Let's split dataset by using function `train_test_split()`. We need to pass 3 parameters features, target, and test\_set size. Additionally, We can use `random_state` to select records randomly.

```
# split X and y into training and testing sets
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.65,random_state=0)
```

Here, the Dataset is broken into two parts in a ratio of 75:25. It means 75% data will be used for model training and 25% for model testing.

**Model Development and Prediction** First, import the Logistic Regression module and create a Logistic Regression classifier object using `LogisticRegression()` function.

Then, fit our model on the train set using `fit()` and perform prediction on the test set using `predict()`.

```
#import the class
from sklearn.linear_model import LogisticRegression

# instantiate the model (using the default parameters)
logreg = LogisticRegression()

# fit the model with data
logreg.fit(X_train,y_train)

#
y_pred=logreg.predict(X_test)

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
`extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,`

## Model Evaluation using Confusion Matrix

A confusion matrix is a table that is used to evaluate the performance of a classification model. We can also visualize the performance of an algorithm. The fundamental of a confusion matrix is the number of correct and incorrect predictions are summed up class-wise.

```
# import the metrics class
from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
cnf_matrix

array([[174,  28],
       [ 56,  43]])
```

Here, we can see the confusion matrix in the form of the array object. The dimension of this matrix is 2\*2 because this model is binary classification. We have two classes 0 and 1. Diagonal values represent accurate predictions, while non-diagonal elements are inaccurate predictions. In the output, 174 and 28 are actual predictions, and 56 and 43 are incorrect predictions.


Visualizing Confusion Matrix using Heatmap Let's visualize the results of the model in the form of a confusion matrix using matplotlib and seaborn.

Here, we will visualize the confusion matrix using Heatmap.

```
# import required modules
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual Chd')
plt.xlabel('Predicted Chd')

Text(0.5, 257.44, 'Predicted Chd')
```



**Confusion Matrix Evaluation Metrics** Let's evaluate the model using model evaluation metrics such as accuracy, precision, and recall.

```
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
print("Precision:", metrics.precision_score(y_test, y_pred))
print("Recall:", metrics.recall_score(y_test, y_pred))
```

```
Accuracy: 0.7209302325581395
Precision: 0.6056338028169014
Recall: 0.43434343434343436
```

Well, we got a classification rate of 72%, considered as good accuracy.


**Precision:** Precision is about being precise, i.e., how accurate our model is. In other words, we can say, when a model makes a prediction, how often it is correct. In our prediction case, when our Logistic Regression model predicted patients are going to suffer from diabetes, that patients have 60% of the time.

**Recall:** If there are patients who have diabetes in the test set and our Logistic Regression model can identify it 43% of the time.

**ROC Curve** Receiver Operating Characteristic(ROC) curve is a plot of the true positive rate against the false positive rate. It shows the tradeoff between sensitivity and specificity.



```
y_pred_proba = logreg.predict_proba(X_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



**AUC score for the case is 0.73. AUC score 1 represents perfect classifier, and 0.5 represents a worthless classifier.**

**Conclusion** In this Notebook we were able to measure and evaluate the Accuracy, Recall, Precision of the data thoroughly.